

Homework 5 – Due Monday October 12, 2015

Cheryl Calhoun

09/29/2015

Assignment #5

Partly inspired by recent adventures, but connected to everyone's experiences when travelling to Gainesville, we will analyze delays associated with airline travel. In particular, the outcome of interest is whether or not a particular flight was noticeably delayed in its arrival to its destination.

The Data Set

Our data set contains the following variables about flights:

- ArrDelay: arrival delay in minutes (negative means an early arrival)
- Month: month of the flight (1 = January; 12 = December)
- DayofWeek: day of the week (1 = Monday; 7 = Sunday)
- AirTime: total time in the air in minutes
- Distance: total distance of the flight in miles
- TaxiIn: time taxiing out to the runway in minutes
- TaxiOut: time taxiing to the gate after landing in minutes
- DepDelay: departure delay in minutes (negative means an early departure)
- Diverted: was the plane diverted? 1 = Yes, 0 = No
- Carrier: was there a delay due to the carrier/company? 1 = Yes, 0 = No
- Weather: was there a delay due to weather? 1 = Yes, 0 = No
- NAS: was there a delay due to the National Air System? 1 = Yes, 0 = No
- Security: was there a delay due to security problems? 1 = Yes, 0 = No
- LateAircraft: was there a delay due to a late aircraft? 1 = Yes, 0 = No

We will begin by loading the data file from its source,

<http://www.acthomas.ca/FSSS/data/airline-data.csv>, using the `read.table` command. We then review the data set to see how many observations (rows) our data set has and confirm that the column names correspond to the variables above.

```
## Load airline data set.
airline <- read.table ("http://www.acthomas.ca/FSSS/data/airline-data.csv", se
p = ",", header=TRUE)

## Review contents of airline data set.
is.data.frame(airline)
```

```
## [1] TRUE
```

```
## Review column (variable) names and data types for each variable.  
str(airline)
```

```
## 'data.frame':    4887 obs. of  15 variables:  
##  $ X          : int  1 2 3 4 5 6 7 8 9 10 ...  
##  $ ArrDelay   : int  -7 57 -15 -4 -5 19 -5 30 106 20 ...  
##  $ Month      : int  2 11 12 3 12 5 12 1 12 3 ...  
##  $ DayOfWeek  : int  4 4 2 3 7 3 5 5 1 1 ...  
##  $ AirTime    : int  244 66 118 73 52 88 111 68 93 177 ...  
##  $ Distance   : int  1814 413 972 444 304 629 758 308 675 1389 ...  
##  $ TaxiIn     : int  8 4 8 3 6 5 6 6 12 4 ...  
##  $ TaxiOut    : int  17 13 12 8 13 12 11 22 19 14 ...  
##  $ DepDelay   : int  0 79 2 -3 -1 33 1 9 97 34 ...  
##  $ Diverted   : int  0 0 0 0 0 0 0 0 0 0 ...  
##  $ Carrier    : int  0 1 0 0 0 1 0 1 1 1 ...  
##  $ Weather    : int  0 0 0 0 0 0 0 0 0 0 ...  
##  $ NAS        : int  0 0 0 0 0 0 0 1 1 0 ...  
##  $ Security   : int  0 0 0 0 0 0 0 0 0 0 ...  
##  $ LateAircraft: int  0 1 0 0 0 0 0 0 0 1 ...
```

```
summary(airline)
```

```
##           X           ArrDelay           Month           DayOfWeek
## Min.      : 1      Min.      : -60.000      Min.      : 1.000      Min.      :1.000
## 1st Qu.:1250      1st Qu.: -10.000      1st Qu.: 4.000      1st Qu.:2.000
## Median :2496      Median :  -2.000      Median : 6.000      Median :4.000
## Mean   :2498      Mean   :  8.456      Mean   : 6.408      Mean   :3.915
## 3rd Qu.:3746      3rd Qu.: 11.000      3rd Qu.: 9.000      3rd Qu.:6.000
## Max.    :5000      Max.    :1092.000      Max.    :12.000      Max.    :7.000
##      AirTime      Distance      TaxiIn      TaxiOut
## Min.      : 9.0      Min.      : 31.0      Min.      : 0.000      Min.      : 2.00
## 1st Qu.: 56.0      1st Qu.: 328.0      1st Qu.: 4.000      1st Qu.: 10.00
## Median : 84.0      Median : 573.0      Median : 6.000      Median : 14.00
## Mean   :102.8      Mean   : 719.1      Mean   : 6.755      Mean   : 16.38
## 3rd Qu.:131.0      3rd Qu.: 944.0      3rd Qu.: 8.000      3rd Qu.: 19.00
## Max.    :580.0      Max.    :4502.0      Max.    :171.000      Max.    :192.00
##      DepDelay      Diverted      Carrier      Weather
## Min.      : -29.00      Min.      :0      Min.      :0.00000      Min.      :0.00000
## 1st Qu.:  -4.00      1st Qu.:0      1st Qu.:0.00000      1st Qu.:0.00000
## Median :  -1.00      Median :0      Median :0.00000      Median :0.00000
## Mean   : 10.21      Mean   :0      Mean   :0.09597      Mean   :0.01514
## 3rd Qu.:  8.00      3rd Qu.:0      3rd Qu.:0.00000      3rd Qu.:0.00000
## Max.    :1099.00      Max.    :0      Max.    :1.00000      Max.    :1.00000
##      NAS      Security      LateAircraft
## Min.      :0.0000      Min.      :0.0000000      Min.      :0.0000
## 1st Qu.:0.0000      1st Qu.:0.0000000      1st Qu.:0.0000
## Median :0.0000      Median :0.0000000      Median :0.0000
## Mean   :0.1306      Mean   :0.0004092      Mean   :0.1076
## 3rd Qu.:0.0000      3rd Qu.:0.0000000      3rd Qu.:0.0000
## Max.    :1.0000      Max.    :1.0000000      Max.    :1.0000
```

Results: The data file, `airline-data.csv`, is a comma delimited file. We used the `read.table()` command with `sep=","` to read the `.csv` file into a data frame. Next we confirmed it was a data frame using `is.data.frame()` which returned a value of `TRUE`. Then we used `str()` to determine there are 4887 observations with 15 total variables, that the column names do in fact correspond to those provided in the assignment, and to determine the data types of each of the variables. We then used `summary()` to get an overview of the descriptive statistics for the variables contained in the data set.

Converting to factors

The first step in our analysis, will be to determine which two numerical columns must be converted to factors and use `mutate()` to do so. We will then confirm the counts of these factors using the `summary` function applied only to those two columns.

```
## Convert variables Month and DayOfWeek to data type factor.
airline.2 <- airline %>%
  mutate (Month=as.factor(Month),
          DayOfWeek=as.factor(DayOfWeek))

## Review data set to verify conversion worked properly.
str(airline.2)
```

```
## 'data.frame':    4887 obs. of  15 variables:
##  $ X              : int  1 2 3 4 5 6 7 8 9 10 ...
##  $ ArrDelay       : int  -7 57 -15 -4 -5 19 -5 30 106 20 ...
##  $ Month          : Factor w/ 12 levels "1","2","3","4",...: 2 11 12 3 12 5 12
##                    1 12 3 ...
##  $ DayOfWeek      : Factor w/ 7 levels "1","2","3","4",...: 4 4 2 3 7 3 5 5 1
##                    1 ...
##  $ AirTime        : int  244 66 118 73 52 88 111 68 93 177 ...
##  $ Distance       : int  1814 413 972 444 304 629 758 308 675 1389 ...
##  $ TaxiIn         : int   8 4 8 3 6 5 6 6 12 4 ...
##  $ TaxiOut        : int  17 13 12 8 13 12 11 22 19 14 ...
##  $ DepDelay       : int   0 79 2 -3 -1 33 1 9 97 34 ...
##  $ Diverted       : int   0 0 0 0 0 0 0 0 0 0 ...
##  $ Carrier        : int   0 1 0 0 0 1 0 1 1 1 ...
##  $ Weather        : int   0 0 0 0 0 0 0 0 0 0 ...
##  $ NAS            : int   0 0 0 0 0 0 0 1 1 0 ...
##  $ Security       : int   0 0 0 0 0 0 0 0 0 0 ...
##  $ LateAircraft   : int   0 1 0 0 0 0 0 0 0 1 ...
```

```
## Use summary to look at each of the converted variables to verify new data looks as is expected.
summary(airline.2$Month)
```

```
##      1      2      3      4      5      6      7      8      9     10     11     12
## 416 381 414 441 428 440 415 425 358 424 355 390
```

```
summary(airline.2$DayOfWeek)
```

```
##      1      2      3      4      5      6      7
## 715 730 736 726 716 580 684
```

Results: Used `as.factor()` to convert Month and DayOfWeek to factors, and `mutate()` to replace the original values for Month and DayOfWeek with the new factor variables. Saved new data.frame as `airline.2`. Used `str()` to verify Month and DayOfWeek are now factors. Used `'summary()'` to view the levels and frequencies of each level.

Creating a variable for delays over 10 minutes

Our objective will be to see if a flight's arrival is delayed by more than 10 minutes. To do this, we will create a new variable in this data frame using `mutate()` corresponding to this True/False outcome.

```
## Calculate number of flights where arrival delay was more than 10 minutes, as  
sign to new variable `Delayed10`.  
airline.3 <- airline.2 %>%  
  mutate(Delayed10 = (ArrDelay > 10))  
  
## Review data set to verify calculation worked correctly.  
str(airline.3)
```

```
## 'data.frame':   4887 obs. of  16 variables:  
##  $ X           : int  1 2 3 4 5 6 7 8 9 10 ...  
##  $ ArrDelay     : int  -7 57 -15 -4 -5 19 -5 30 106 20 ...  
##  $ Month        : Factor w/ 12 levels "1","2","3","4",...: 2 11 12 3 12 5 12  
1 12 3 ...  
##  $ DayOfWeek    : Factor w/ 7 levels "1","2","3","4",...: 4 4 2 3 7 3 5 5 1  
1 ...  
##  $ AirTime      : int  244 66 118 73 52 88 111 68 93 177 ...  
##  $ Distance     : int  1814 413 972 444 304 629 758 308 675 1389 ...  
##  $ TaxiIn       : int  8 4 8 3 6 5 6 6 12 4 ...  
##  $ TaxiOut      : int  17 13 12 8 13 12 11 22 19 14 ...  
##  $ DepDelay     : int  0 79 2 -3 -1 33 1 9 97 34 ...  
##  $ Diverted     : int  0 0 0 0 0 0 0 0 0 0 ...  
##  $ Carrier      : int  0 1 0 0 0 1 0 1 1 1 ...  
##  $ Weather      : int  0 0 0 0 0 0 0 0 0 0 ...  
##  $ NAS          : int  0 0 0 0 0 0 0 1 1 0 ...  
##  $ Security     : int  0 0 0 0 0 0 0 0 0 0 ...  
##  $ LateAircraft : int  0 1 0 0 0 0 0 0 0 1 ...  
##  $ Delayed10    : logi  FALSE TRUE FALSE FALSE FALSE TRUE ...
```

```
summary(airline.3$Delayed10)
```

```
##      Mode   FALSE    TRUE   NA's  
## logical   3606    1281      0
```

Results: Created new variable `Delayed10` to indicate whether or not a flight's arrival is delayed by more than 10 minutes. Saved new data.frame as `airline.3`. Used `str()` to verify the new variable has been created, and view the resulting data type. Used `summary()` to view descriptive statistics and verify the new variable was calculated correctly. `Delayed10` now contains logical data indicating true for flights with delays > 10 minutes, and false for flights with delays <= 10 minutes.

Fitting the binary linear model

Now we will use `glm()` to fit a binary linear model to the Delayed10 outcome. We will include all variables initially, and then look at whether or not any variables can be omitted if it is determined they have no bearing on the analysis.

```
## Use `glm` to fit a binary linear model to `Delayed10` with all variables in  
the data set  
delayed.model <- glm (Delayed10 ~ Month + DayOfWeek + AirTime + Distance + Taxi  
In + TaxiOut + DepDelay + Carrier + Weather + NAS + Security + LateAircraft -  
1, data=airline.3, family=binomial)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## View resulting model and summary statistics.  
summary(delayed.model)
```

```
##
## Call:
## glm(formula = Delayed10 ~ Month + DayOfWeek + AirTime + Distance +
##       TaxiIn + TaxiOut + DepDelay + Carrier + Weather + NAS + Security +
##       LateAircraft - 1, family = binomial, data = airline.3)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.7468  -0.2281  -0.1311   0.0000   3.3459
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## Month1         -7.045e+00  4.867e-01 -14.475 < 2e-16 ***
## Month2         -6.944e+00  4.992e-01 -13.909 < 2e-16 ***
## Month3         -7.295e+00  5.194e-01 -14.045 < 2e-16 ***
## Month4         -6.440e+00  4.454e-01 -14.458 < 2e-16 ***
## Month5         -6.252e+00  4.298e-01 -14.545 < 2e-16 ***
## Month6         -6.683e+00  4.583e-01 -14.583 < 2e-16 ***
## Month7         -7.120e+00  4.957e-01 -14.363 < 2e-16 ***
## Month8         -7.029e+00  4.923e-01 -14.276 < 2e-16 ***
## Month9         -6.889e+00  4.998e-01 -13.782 < 2e-16 ***
## Month10        -7.053e+00  4.855e-01 -14.528 < 2e-16 ***
## Month11        -6.997e+00  5.104e-01 -13.708 < 2e-16 ***
## Month12        -6.974e+00  4.822e-01 -14.464 < 2e-16 ***
## DayOfWeek2      4.340e-01  3.135e-01  1.384  0.1663
## DayOfWeek3      3.467e-01  3.157e-01  1.098  0.2721
## DayOfWeek4      4.061e-01  3.112e-01  1.305  0.1919
## DayOfWeek5     -5.248e-02  3.395e-01 -0.155  0.8771
## DayOfWeek6      5.465e-01  3.317e-01  1.648  0.0994 .
## DayOfWeek7      3.349e-01  3.218e-01  1.041  0.2980
## AirTime         5.354e-02  7.866e-03  6.806 1.00e-11 ***
## Distance        -6.744e-03  9.918e-04 -6.799 1.05e-11 ***
## TaxiIn          1.072e-01  1.710e-02  6.270 3.62e-10 ***
## TaxiOut          1.087e-01  9.977e-03  10.895 < 2e-16 ***
## DepDelay        1.794e-01  1.061e-02  16.903 < 2e-16 ***
## Carrier         1.949e+01  7.112e+02  0.027  0.9781
## Weather         1.659e+01  1.734e+03  0.010  0.9924
## NAS             2.159e+01  6.912e+02  0.031  0.9751
## Security        -1.125e+01  1.725e+04 -0.001  0.9995
## LateAircraft    1.774e+01  6.866e+02  0.026  0.9794
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6774.8  on 4887  degrees of freedom
## Residual deviance: 1096.9  on 4859  degrees of freedom
## AIC: 1152.9
```

```
##  
## Number of Fisher Scoring iterations: 20
```

Results: The resulting model shows 6 variables which are statistically significant. One of these is `Month` which appears as all 12 months. `DayOfWeek` is not producing statistically significant results. `Diverted` has been eliminated because it does not seem to contain any significant data as all observations = 0.

Note: `glm` creates a warning “## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred” when run for “family=binomial”. This error is generated due to a condition where one or more of the fitted probabilities are extremely close to zero or one. This is not a fatal flaw but should be noted.

Standardizing the variables

Now let's look at the standardized model to see if we can improve the fit. We will create a new data frame that standardizes the variables in question by subtract their mean and divide by their standard deviation. We will use `qplot` to look at the resulting standardized variables to confirm that they appear to have the correct distributions.

```
## Convert AirTime, Distance, TaxiIn, TaxiOut and DepDelay to standardized values.  
airline.zs <- airline.3 %>%  
  mutate(AirTime = (AirTime - mean(AirTime)) / sd(AirTime)) %>%  
  mutate(Distance = (Distance - mean(Distance)) / sd(Distance)) %>%  
  mutate(TaxiIn = (TaxiIn - mean(TaxiIn)) / sd(TaxiIn)) %>%  
  mutate(TaxiOut = (TaxiOut - mean(TaxiOut)) / sd(TaxiOut)) %>%  
  mutate(DepDelay = (DepDelay - mean(DepDelay)) / sd(DepDelay))  
  
# View resulting data set and summary statistics to verify standardized values appear correct.  
str(airline.zs)
```

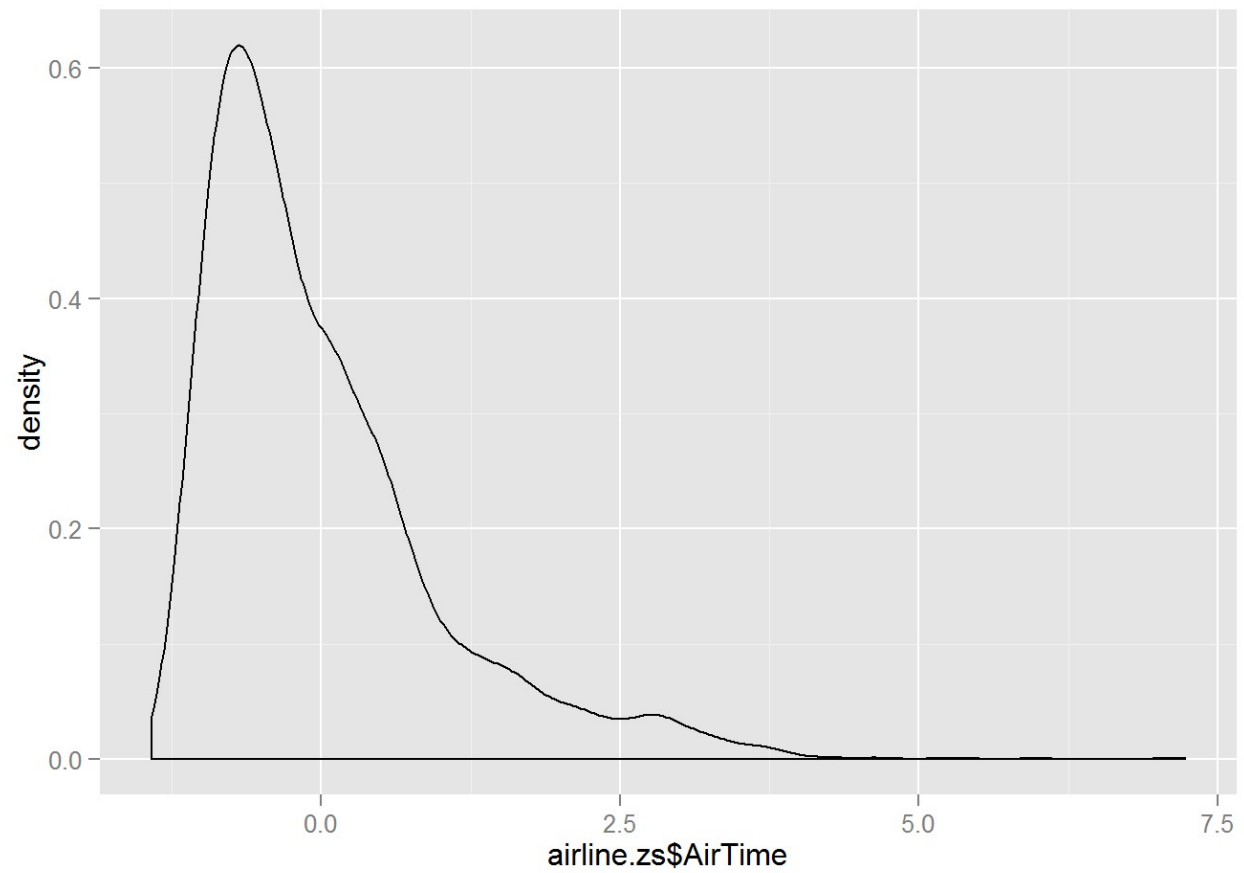


```
## 'data.frame':    4887 obs. of  16 variables:
##  $ X              : int  1 2 3 4 5 6 7 8 9 10 ...
##  $ ArrDelay       : int  -7 57 -15 -4 -5 19 -5 30 106 20 ...
##  $ Month          : Factor w/ 12 levels "1","2","3","4",...: 2 11 12 3 12 5 12
1 12 3 ...
##  $ DayOfWeek      : Factor w/ 7 levels "1","2","3","4",...: 4 4 2 3 7 3 5 5 1
1 ...
##  $ AirTime        : num  2.139 -0.557 0.23 -0.451 -0.769 ...
##  $ Distance       : num  1.991 -0.557 0.46 -0.5 -0.755 ...
##  $ TaxiIn         : num  0.232 -0.514 0.232 -0.7 -0.141 ...
##  $ TaxiOut        : num  0.0557 -0.3009 -0.39 -0.7465 -0.3009 ...
##  $ DepDelay       : num  -0.277 1.868 -0.223 -0.359 -0.305 ...
##  $ Diverted       : int  0 0 0 0 0 0 0 0 0 0 ...
##  $ Carrier        : int  0 1 0 0 0 1 0 1 1 1 ...
##  $ Weather        : int  0 0 0 0 0 0 0 0 0 0 ...
##  $ NAS            : int  0 0 0 0 0 0 0 1 1 0 ...
##  $ Security       : int  0 0 0 0 0 0 0 0 0 0 ...
##  $ LateAircraft   : int  0 1 0 0 0 0 0 0 0 1 ...
##  $ Delayed10      : logi  FALSE TRUE FALSE FALSE FALSE TRUE ...
```


```
summary(airline.zs)
```

```
##           X           ArrDelay           Month           DayOfWeek
## Min.      : 1      Min.      : -60.000      4           : 441      1:715
## 1st Qu.:1250      1st Qu.: -10.000      6           : 440      2:730
## Median :2496      Median :   -2.000      5           : 428      3:736
## Mean    :2498      Mean     :    8.456      8           : 425      4:726
## 3rd Qu.:3746      3rd Qu.:  11.000     10           : 424      5:716
## Max.    :5000      Max.     :1092.000      1           : 416      6:580
##                                     (Other):2313      7:684
##           AirTime           Distance           TaxiIn           TaxiOut
## Min.      :-1.4206      Min.      :-1.2513      Min.      :-1.2601      Min.      :-1.2814
## 1st Qu.: -0.7087      1st Qu.: -0.7112      1st Qu.: -0.5139      1st Qu.: -0.5683
## Median : -0.2845      Median : -0.2657      Median : -0.1408      Median : -0.2117
## Mean     : 0.0000      Mean     : 0.0000      Mean     : 0.0000      Mean     : 0.0000
## 3rd Qu.: 0.4274      3rd Qu.: 0.4090      3rd Qu.: 0.2323      3rd Qu.: 0.2340
## Max.     : 7.2283      Max.     : 6.8793      Max.     :30.6399      Max.     :15.6545
##
##           DepDelay           Diverted           Carrier           Weather
## Min.      :-1.06495      Min.      :0      Min.      :0.00000      Min.      :0.00000
## 1st Qu.: -0.38598      1st Qu.:0      1st Qu.:0.00000      1st Qu.:0.00000
## Median : -0.30451      Median :0      Median :0.00000      Median :0.00000
## Mean     : 0.00000      Mean     :0      Mean     :0.09597      Mean     :0.01514
## 3rd Qu.: -0.06008      3rd Qu.:0      3rd Qu.:0.00000      3rd Qu.:0.00000
## Max.     :29.56996      Max.      :0      Max.     :1.00000      Max.     :1.00000
##
##           NAS           Security           LateAircraft           Delayed10
## Min.      :0.0000      Min.      :0.0000000      Min.      :0.0000      Mode :logical
## 1st Qu.:0.0000      1st Qu.:0.0000000      1st Qu.:0.0000      FALSE:3606
## Median :0.0000      Median :0.0000000      Median :0.0000      TRUE :1281
## Mean     :0.1306      Mean     :0.0004092      Mean     :0.1076      NA's :0
## 3rd Qu.:0.0000      3rd Qu.:0.0000000      3rd Qu.:0.0000
## Max.     :1.0000      Max.     :1.0000000      Max.     :1.0000
##
```

```
## Plot two of the standardized variable to verify they are in the range of sta
ndardized scores.
qplot(airline.zs$AirTime, geom="density")
```



```
qplot(airline.zs$Distance, geom="density")
```



Results: We now have a new data frame, `airline.zs`, which contains the standardized z scores for each non 0/1 predictor. The resulting z scores produce plots with distributions clustered around “0”, verifying the z scores were correctly calculated.

Fitting the binary linear model with standardized data

Again, we will use `glm()` to fit a binary linear model to the `Delayed10` outcome using the new data set which contains standardized variables for predictors.. We will include all variables initially, and then look at whether or not any variables can be omitted if it is determined they have no bearing on the analysis. but with these standardized variables for predictors instead. Which variables have the greatest effect size in each regression?

```
## Use `glm` to fit a binary linear model to Delayed10 using the standardized v  
ariables in the data set  
delay.model.zs <- glm (Delayed10 ~ Month + DayOfWeek + AirTime + Distance + Tax  
iIn + TaxiOut + DepDelay + Carrier + Weather + NAS + Security + LateAircraft  
-1, data=airline.zs, family=binomial)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## View resulting model and summary statistics.  
##delay.model.zs  
summary(delay.model.zs)
```

```
##
## Call:
## glm(formula = Delayed10 ~ Month + DayOfWeek + AirTime + Distance +
##       TaxiIn + TaxiOut + DepDelay + Carrier + Weather + NAS + Security +
##       LateAircraft - 1, family = binomial, data = airline.zs)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.7468  -0.2281  -0.1311   0.0000   3.3459
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## Month1         -2.055e+00  3.673e-01  -5.596 2.20e-08 ***
## Month2         -1.955e+00  3.785e-01  -5.164 2.42e-07 ***
## Month3         -2.306e+00  4.034e-01  -5.716 1.09e-08 ***
## Month4         -1.450e+00  3.388e-01  -4.279 1.87e-05 ***
## Month5         -1.263e+00  3.157e-01  -4.000 6.34e-05 ***
## Month6         -1.694e+00  3.364e-01  -5.035 4.79e-07 ***
## Month7         -2.130e+00  3.704e-01  -5.751 8.87e-09 ***
## Month8         -2.039e+00  3.864e-01  -5.277 1.31e-07 ***
## Month9         -1.899e+00  3.845e-01  -4.939 7.84e-07 ***
## Month10        -2.063e+00  3.768e-01  -5.476 4.34e-08 ***
## Month11        -2.007e+00  4.092e-01  -4.905 9.32e-07 ***
## Month12        -1.985e+00  3.674e-01  -5.402 6.60e-08 ***
## DayOfWeek2      4.340e-01  3.135e-01   1.384  0.1663
## DayOfWeek3      3.467e-01  3.157e-01   1.098  0.2721
## DayOfWeek4      4.061e-01  3.112e-01   1.305  0.1919
## DayOfWeek5     -5.248e-02  3.394e-01  -0.155  0.8771
## DayOfWeek6      5.465e-01  3.317e-01   1.648  0.0994 .
## DayOfWeek7      3.349e-01  3.218e-01   1.041  0.2980
## AirTime         3.535e+00  5.193e-01   6.806 1.00e-11 ***
## Distance        -3.708e+00  5.454e-01  -6.799 1.05e-11 ***
## TaxiIn          5.747e-01  9.166e-02   6.270 3.62e-10 ***
## TaxiOut         1.219e+00  1.119e-01  10.895 < 2e-16 ***
## DepDelay        6.605e+00  3.908e-01  16.903 < 2e-16 ***
## Carrier         1.949e+01  7.112e+02   0.027  0.9781
## Weather         1.659e+01  1.734e+03   0.010  0.9924
## NAS             2.159e+01  6.912e+02   0.031  0.9751
## Security        -1.125e+01  1.725e+04  -0.001  0.9995
## LateAircraft    1.774e+01  6.866e+02   0.026  0.9794
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6774.8  on 4887  degrees of freedom
## Residual deviance: 1096.9  on 4859  degrees of freedom
## AIC: 1152.9
```

```
##  
## Number of Fisher Scoring iterations: 20
```

Results: The resulting model shows 6 variables which are statistically significant. One of these is `Month` which appears as all 12 months. `DayOfWeek` is not producing statistically significant results. `Diverted` has been eliminated because it does not seem to contain any significant data as all observations = 0. `DepDelay` has the greatest effect size in both models.

Note: `glm` creates a warning “## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred” when run for “family=binomial”. This error is generated due to a condition where one or more of the fitted probabilities are extremely close to zero or one. This is not a fatal flaw but should be noted.

Creating the design/predictors matrix

Now we will produce the design/predictors matrix using the `model.matrix()` function. We will verify the number of columns corresponds to the number of coefficients in your previous `glm()` output.

```
## Produce the design/predictors matrix for airline.delays.zs.  
delay.matrix <- model.matrix(~ 0 + Month + DayOfWeek + AirTime + Distance + TaxiIn + TaxiOut + DepDelay + Carrier + Weather + NAS + Security + LateAircraft, airline.zs)  
  
## Verify the number of columns in airline.delay.matrix corresponds to the number of coefficients in the previous `glm()` output.  
length(delay.model.zs$coefficients)
```

```
## [1] 28
```

```
dim(delay.matrix)[2]
```


```
## [1] 28
```

Results: The `delay.matrix` has the same number of columns (28) as the number of coefficients in the previous ‘`delay.model.zs`’ (28).

Lasso effect

We will use `cv.glmnet()` with the Lasso (`alpha=1`) to run a penalized linear model for quality as the outcome with all predictors as previously done, for this data frame.

```
##Use cv.glmnet() with Lasso (`alpha=1`) to run the penalized linear model for
delayed.10.
## First Validation Step
delay.model.cv = cv.glmnet (delay.matrix, airline.zs$Delayed10, alpha=1 )
picked <- which (delay.model.cv$lambda == delay.model.cv$lambda.min)
plot(delay.model.cv)
```



Results: This plot shows the cross-validation mean squared error (MSE) as a function of log(lambda) curve (red dotted line), including the upper and lower standard deviation curves. The dotted lines represent lambda.min and lambda.min plus one standard error. As lambda gets smaller, the curve flattens out. The numbers across the top of the plot indicate how many non-zero predictors are in the model at each level of lambda.

```
## The `lambda` that produces the smallest cross-validated error is:
delay.model.cv$lambda.min
```

```
## [1] 0.0005114175
```

```
## The cross validated error is:
delay.model.cv$cvm[picked]
```



```
## [1] 0.06210824
```

```
## How much of a reduction is this in cross-validated error from the basic fit  
model in Question 3?  
summary(delay.model.zs)
```

```
##
## Call:
## glm(formula = Delayed10 ~ Month + DayOfWeek + AirTime + Distance +
##       TaxiIn + TaxiOut + DepDelay + Carrier + Weather + NAS + Security +
##       LateAircraft - 1, family = binomial, data = airline.zs)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.7468  -0.2281  -0.1311   0.0000   3.3459
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## Month1         -2.055e+00  3.673e-01  -5.596 2.20e-08 ***
## Month2         -1.955e+00  3.785e-01  -5.164 2.42e-07 ***
## Month3         -2.306e+00  4.034e-01  -5.716 1.09e-08 ***
## Month4         -1.450e+00  3.388e-01  -4.279 1.87e-05 ***
## Month5         -1.263e+00  3.157e-01  -4.000 6.34e-05 ***
## Month6         -1.694e+00  3.364e-01  -5.035 4.79e-07 ***
## Month7         -2.130e+00  3.704e-01  -5.751 8.87e-09 ***
## Month8         -2.039e+00  3.864e-01  -5.277 1.31e-07 ***
## Month9         -1.899e+00  3.845e-01  -4.939 7.84e-07 ***
## Month10        -2.063e+00  3.768e-01  -5.476 4.34e-08 ***
## Month11        -2.007e+00  4.092e-01  -4.905 9.32e-07 ***
## Month12        -1.985e+00  3.674e-01  -5.402 6.60e-08 ***
## DayOfWeek2      4.340e-01  3.135e-01   1.384  0.1663
## DayOfWeek3      3.467e-01  3.157e-01   1.098  0.2721
## DayOfWeek4      4.061e-01  3.112e-01   1.305  0.1919
## DayOfWeek5     -5.248e-02  3.394e-01  -0.155  0.8771
## DayOfWeek6      5.465e-01  3.317e-01   1.648  0.0994 .
## DayOfWeek7      3.349e-01  3.218e-01   1.041  0.2980
## AirTime         3.535e+00  5.193e-01   6.806 1.00e-11 ***
## Distance        -3.708e+00  5.454e-01  -6.799 1.05e-11 ***
## TaxiIn          5.747e-01  9.166e-02   6.270 3.62e-10 ***
## TaxiOut         1.219e+00  1.119e-01  10.895 < 2e-16 ***
## DepDelay        6.605e+00  3.908e-01  16.903 < 2e-16 ***
## Carrier         1.949e+01  7.112e+02   0.027  0.9781
## Weather         1.659e+01  1.734e+03   0.010  0.9924
## NAS             2.159e+01  6.912e+02   0.031  0.9751
## Security        -1.125e+01  1.725e+04  -0.001  0.9995
## LateAircraft    1.774e+01  6.866e+02   0.026  0.9794
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6774.8  on 4887  degrees of freedom
## Residual deviance: 1096.9  on 4859  degrees of freedom
## AIC: 1152.9
```


```
##  
## Number of Fisher Scoring iterations: 20
```

Results: The `lambda` that produces the smallest value is [69] = 5.114175310^{-4} . The cross-validated error for `red.quality.cv` [69] is 0.0621082.

Ridge penalty

Let's run `cv.glmnet()` again, but now using the Ridge penalty (`alpha=0`).

```
##Use cv.glmnet() with Lasso (`alpha=1`) to run the penalized linear model for  
delayed.10.  
## First Validation Step  
delay.model.cv2 = cv.glmnet (delay.matrix, airline.zs$Delayed10, alpha=0 )  
picked2 <- which (delay.model.cv2$lambda == delay.model.cv2$lambda.min)  
plot(delay.model.cv2)
```



Results: This plot shows the cross-validation mean squared error (MSE) as a function of $\log(\lambda)$ curve (red dotted line), including the upper and lower standard deviation curves. The dotted lines represent λ_{\min} and λ_{\min} plus one standard error. As λ gets smaller, the curve flattens out. The numbers across the top of the plot indicate how many non-zero predictors are in the model at each level of λ .

```
## The `lambda` that produces the smallest cross-validated error is:  
delay.model.cv2$lambda.min
```

```
## [1] 0.03138011
```

```
## The cross validated error is:  
delay.model.cv2$cvm[picked2]
```

```
## [1] 0.06236552
```


```
## How much of a reduction is this in cross-validated error from the basic fit  
model in Question 3?  
## summary(delay.model.zs)  
## delay.model.zs
```

Results: The λ that produces the smallest value is [99] = 0.0313801. The cross-validated error for red.quality.cv [99] is NA.

Balancing both penalties

Let's run `cv.glmnet` one more time, but now using an even combination of the two penalties ($\alpha=0.5$).

```
## Use cv.glmnet() with Lasso (`alpha=1`) to run the penalized linear model for  
delayed.10.  
## First Validation Step  
delay.model.cv3 = cv.glmnet (delay.matrix, airline.zs$Delayed10, alpha=0.5 )  
picked3 <- which (delay.model.cv3$lambda == delay.model.cv3$lambda.min)  
plot(delay.model.cv3)
```



Results: This plot shows the cross-validation mean squared error (MSE) as a function of $\log(\lambda)$ curve (red dotted line), including the upper and lower standard deviation curves. The dotted lines represent λ_{\min} and λ_{\min} plus one standard error. As λ gets smaller, the curve flattens out. The numbers across the top of the plot indicate how many non-zero predictors are in the model at each level of λ .

```
## The `lambda` that produces the smallest cross-validated error is:
delay.model.cv3$lambda.min
```

```
## [1] 0.0008491756
```

```
## The cross validated error is:
delay.model.cv3$cvm[picked3]
```

```
## [1] 0.06199374
```

```
## How much of a reduction is this in cross-validated error from the basic fit
model in Question 3?
## summary(delay.model.zs)
## delay.model.zs
```

Results: The `lambda` that produces the smallest value is `[71] = 5.114175310^{-4}`. The cross-validated error for `red.quality.cv [71]` is 0.0621096. The output from the `glm` generated model does not show adjusted r-squared for comparison.

Finding the smallest lambda

Now let's look at each of our models to determine which one contains the smallest value of `cvm` for the `lambda` we selected.

```
delay.model.cv$cvm[picked]
```

```
## [1] 0.06210824
```

```
delay.model.cv2$cvm[picked2]
```

```
## [1] 0.06236552
```

```
delay.model.cv3$cvm[picked3]
```


```
## [1] 0.06199374
```

Results: The smallest `cvm` value comes from `delay.model.cv3[71]` at 0.0619937.

Fitting the “shrinkage model”

We will use `glmnet()` to fit the “shrinkage” model to this data set using `delay.model.cv3$lambda` which gave us the smallest `cvm`. We will use the same `lambda` series as outputted in the previous steps for each model to look at the `beta` matrix that corresponds to this ideal `lambda`.

```
airline.delay.model = glmnet (delay.matrix, airline.zs$Delayed10, lambda = delay.model.cv3$lambda)
plot(airline.delay.model, label = TRUE)
```



Results: Each curve in the plot represents a coefficient in the model. As the norm increases the coefficients deviate from 0. The numeric scale across the top indicates how many coefficients have deviated from 0.

```
airline.delay.model$beta[,picked3]
```

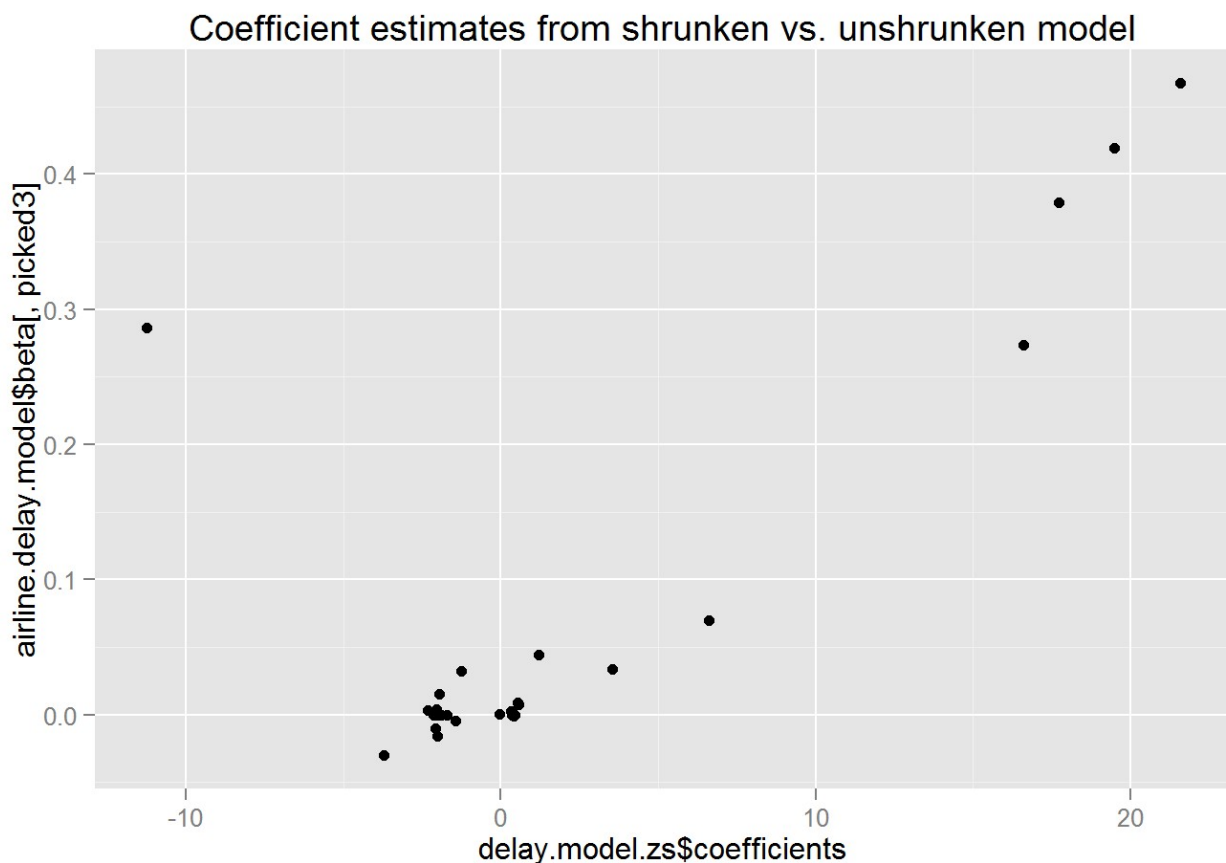
```
##      Month1      Month2      Month3      Month4      Month5
## 0.0037187127 0.0156473331 0.0030852216 -0.0047970955 0.0324174794
##      Month6      Month7      Month8      Month9      Month10
## 0.0000000000 0.0000000000 0.0000000000 0.0000000000 -0.0103927007
##      Month11      Month12      DayOfWeek2      DayOfWeek3      DayOfWeek4
## -0.0159285624 0.0000000000 0.0000000000 0.0000000000 -0.0012428638
##      DayOfWeek5      DayOfWeek6      DayOfWeek7      AirTime      Distance
## 0.0005185574 0.0090237269 0.0027634698 0.0338186086 -0.0299654600
##      TaxiIn      TaxiOut      DepDelay      Carrier      Weather
## 0.0076587143 0.0441382760 0.0699790545 0.4195586937 0.2736560054
##      NAS      Security      LateAircraft
## 0.4673484890 0.2864899728 0.3793499846
```

Results: In this step we calculated the model which corresponds to the minimum `lambda` identified in step 7. The resulting plot and betas are shown. The estimates of `airline.delay.model$beta` for `Month6`, `Month7`, `Month8`, `Month9`, `Month12`, `DaysOfWeeks2`, and `DaysOfWeeks3` have shrunk to zero.

Comparing the unshrunk model to the shrunk model

Finally we will plot the coefficient estimates from the unshrunk models (`delay.model.zs`) compared to the ideal shrunk model (`airline.delay.model`),to demonstrate whether this shrunk estimation produced a noticeably different response.

```
qplot (delay.model.zs$coefficients, airline.delay.model$beta[,picked3], main="Coefficient estimates from shrunk vs. unshrunk model")
```



Results: In this step we plotted the original coefficient estimates from `delay.model.zs` to the coefficient estimates from the shrunk model `airline.delay.model$beta[,picked3]`. There appear to be some differences in coefficients. If there were no differences, the plot would create a straight line.